

# Real-Time Systems

## Resource Access Protocols

# Problems: Priority Inversion


## Assumptions:

- Jobs use resources in a mutually exclusive manner
- Preemptive priority-driven scheduling
- Fixed task priorities
- 1 processor

Declaration  
for the following slides

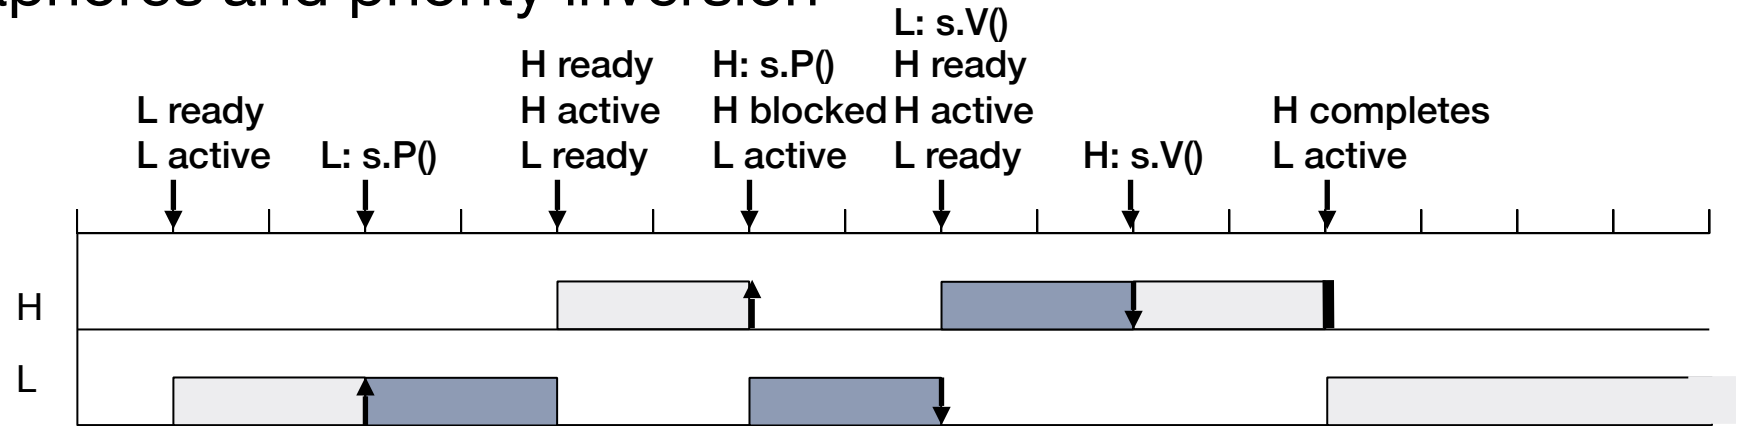
L(R) ↑    U(R) ↓

## Busy waiting and priority inversion

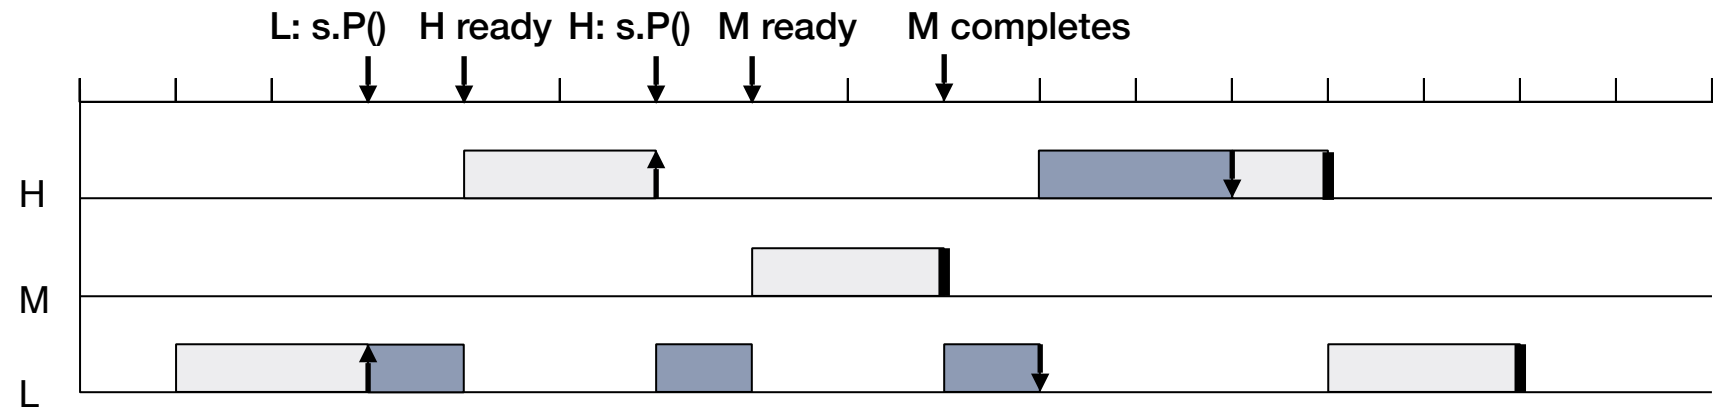


# Problems: Priority Inversion


## Semaphores and priority inversion




M: medium-prioritized job (not using  $s$ )



# Problems: Timing Anomalies



Reduction of resource usage of  $J_3$  by 1.5:



# Problems: Deadlocks

- exclusive resources
- non-preemptive resources
- sequential acquire
- cyclic wait-condition

# Assumptions and Notations

1 processor, preemptive priority-driven scheduling,  
jobs are not self-suspending

- $R_1, \dots, R_r$  resources; nonpreemptable, exclusive
- $L(R_k), U(R_k)$  acquire/release of  $R_k$ ; release: LIFO  
 $\uparrow R_k \quad \downarrow R_k$
- $J_1, \dots, J_n$  jobs
- $J_h, J_l$  job of high/low priority
- $p_1, \dots, p_n$  assigned priorities (highest priority: 1);  
w.l.o.g.:  $J_i$  ordered according to priorities
- $p_i(t)$  current priority of  $J_i$

# Assumptions and Notations

- **Jobs conflict with one another**  
operate with a common resource
- **Jobs contend for a resource**  
one job requests the resource that another job already owns
- **Blocked job**  
scheduler does not grant the requested resource
- **Priority inversion**  
 $J_l$  executes while  $J_h$  is blocked

# Priority Inheritance Protocol

for preemptive priority-driven scheduling

Sha et al., 1990


## Basic Priority-Inheritance Protocol

- (1) **Scheduling Rule**
  - A ready job  $J$  is scheduled according to its current priority  $p(t)$ ; at release time  $t$ :  $p(t) := p$ .
- (2) **Allocation Rule**
  - $J$  requests  $R$  at time  $t$ .
    - (a)  $R$  free:  $R$  is allocated to  $J$  until  $J$  releases  $R$ .
    - (b)  $R$  not free: request is denied,  $J$  is blocked.
- (3) **Priority-Inheritance Rule**
  - When  $J$  becomes blocked by  $J_i$ , then  $J_i$  inherits the current priority of  $J$ , i.e.  $p_i(t) := p(t)$ .
  - $J_i$  executes at this priority until it releases  $R$  at time  $t''$ .
  - Now the priority of  $J_i$  returns to its previous priority:  
 $p_i(t'') := p_i(t')$       $t'$ : time when  $J_i$  acquires  $R$ .



# Priority Inheritance – Example

- 2 jobs: no effect!
- 3 jobs:




# Priority Inheritance – Properties

## Properties


- Priority inheritance is transitive.
- No unbounded uncontrolled priority inversion.
- Priority inheritance does not reduce the blocking times as small as possible.

# Priority Inheritance – Properties



# Priority Inheritance – Properties

Priority inheritance does not prevent deadlocks.



# Priority Ceilings – Notations

Sha/Rajkumar/Lehoczky, 1990

- **Assumptions and Notations**

- 1 processor, preemptive priority-driven scheduling  
no self-suspension
- Assigned priorities  $p_i$  are fixed  
priorities: natural numbers, 1 highest,  $\Omega$  lowest priority
- The resources required by all jobs are known a priori
- **$P(R)$**  priority ceiling of  $R$   
highest priority of all jobs that require  $R$
- **$\hat{P}(t)$**  priority ceiling of the system at time  $t$   
highest priority ceiling of all resources that are in use at time  $t$

# Basic Priority-Ceiling Protocol

- (1) **Scheduling Rule**

- At release time  $t^{rel}$  of  $J$ :  $p(t^{rel}) := p$

- (2) **Allocation Rule**

- $J$  requests  $R$  at time  $t$
- (a)  $R$  held by another job: request denied,  $J$  blocks (“on  $R$ ”)
- (b)  $R$  free:
  - (α)  $p(t) > \hat{P}(t)$ :  $R$  is allocated to  $J$
  - (β) otherwise:  $R$  is allocated to  $J$  only if  $J$  is the job holding the resource(s)  $R'$  with  $P(R') = \hat{P}(t)$ , otherwise  $J$  blocks


- (3) **Priority-Inheritance Rule**

- When  $J$  becomes blocked by  $J_l$ ,  $J_l$  inherits  $J$ 's current priority  $p(t)$
- $J_l$  (preemptively) executes at this priority until it releases every resource whose priority ceiling is at least  $p(t)$
- At that time,  $J_l$ 's priority returns to  $p_l(t')$   
( $t'$ : time when it was granted the resource)


# Basic Priority-Ceiling Protocol – Properties

- Difference to priority inheritance: three ways to blocking


- direct:



- inheritance:





- ceilings:



- Deadlocks can never occur
- There can be no transitive blocking

# Basic Priority-Ceiling Protocol – Example

- A job can be blocked for at most one resource request
- Computation of blocking time – Example:





# Stack-Based Priority-Ceiling Protocol

- **Further Assumptions**

- Common run-time stack for all jobs (no self-suspension)
- Stack space of an active job is on the top of the stack (preemption)
- Stack space is freed when the job completes

- **Protocol**

- (0)  $\hat{P}(t) = \Omega$ , when all  $R$  are free,  
 $\hat{P}(t)$  is updated whenever a resource is allocated or freed
- (1) **Scheduling Rule**
  - After  $J$  is released, it is blocked until  $p > \hat{P}(t)$
  - Priority-driven scheduling based on assigned priorities (!)
- (2) **Allocation Rule**
  - Whenever a job requests a resource, it is granted the resource (!)

- **Properties**

- When a job begins execution, all resources it will ever need are free
- Both protocols result in the same longest blocking time of a job
- Deadlocks cannot occur